

RAJALAKSHMI ENGINEERING COLLEGE

(An Autonomous Institution)

RAJALAKSHMI NAGAR, THANDALAM- 602 105



RAJALAKSHMI ENGINEERING COLLEGE

CS19P18 - DEEP LEARNING CONCEPTS

LABORATORY RECORD NOTEBOOK

NAME: JOEL SUNDARSINGH A

YEAR/SEMESTER: IV/VII

BRANCH: COMPUTER SCIENCE AND ENGINEERING

REGISTER NO:220701110

COLLEGE ROLL NO:2116220701110

ACADEMIC YEAR: 2025 -2026



RAJALAKSHMI ENGINEERING COLLEGE

(An Autonomous Institution)

RAJALAKSHMI NAGAR, THANDALAM- 602 105

BONAFIDE CERTIFICATE

NAME: JOEL SUNDARSINGH A **BRANCH/SECTION:** COMPUTER SCIENCE AND ENGINEERING/B **ACADEMIC YEAR:** 2025 -2026
SEMESTER: VII

REGISTER NO: 220701110

Certified that this is a Bonafide record of work done by the above student in the **CS19P18 - DEEP LEARNING CONCEPTS** during the year 2025 - 2026

Signature of Faculty In-charge

Submitted for the Practical Examination Held on: 25/11/2025

Internal Examiner

External Examiner

INDEX

EX.NO	DATE	NAME OF THE EXPERIMENT	PAGE NO	STAFF SIGN
1	14/07/2025	Create a neural network to recognize handwritten digits using MNIST dataset	5	
2	21/07/2025	Build a Convolutional Neural Network with Keras/TensorFlow	8	
3	28/07/2025	Image Classification on CIFAR-10 Dataset using CNN	11	
4	04/08/2025	Transfer learning with CNN and Visualization	13	
5	25/08/2025	Build a Recurrent Neural Network using Keras/Tensorflow	17	
6	15/09/2025	Sentiment Classification of Text using RNN	19	
7	22/09/2025	Build autoencoders with keras/tensorflow	21	
8	29/09/2025	Object detection with yolo3	24	
9	29/09/2025	Build GAN with Keras/TensorFlow	27	
10	06/10/2025	Mini Project	31	

INSTALLATION AND CONFIGURATION OF TENSORFLOW

Aim:

To install and configure TensorFlow in anaconda environment in Windows 10.

Procedure:

1. Download Anaconda Navigator and install.
2. Open Anaconda prompt
3. Create a new environment dlc with python 3.7 using the following command:
`conda create -n dlc python=3.7`
4. Activate newly created environment dlc using the following command:
`conda activate dlc`
5. In dlc prompt, install tensorflow using the following command:
`pip install tensorflow`
6. Next install Tensorflow-datasets using the following command:
`pip install tensorflow-datasets`
7. Install scikit-learn package using the following command:
`pip install scikit-learn`
8. Install pandas package using the following command:
`pip install pandas`
9. Lastly, install jupyter notebook
`pip install jupyter notebook`
10. Open jupyter notebook by typing the following in dlc prompt:
`jupyter notebook`
11. Click create new and then choose python 3 (ipykernel)
12. Give the name to the file
13. Type the code and click Run button to execute (eg. Type `import tensorflow` and then run)

**EX NO: 1 CREATE A NEURAL NETWORK TO RECOGNIZE HANDWRITTEN
DATE:14/07/2025 DIGITS USING MNIST DATASET**

Aim:

To build a handwritten digit's recognition with MNIST dataset.

Procedure:

1. Download and load the MNIST dataset.
2. Perform analysis and preprocessing of the dataset.
3. Build a simple neural network model using Keras/TensorFlow.
4. Compile and fit the model.
5. Perform prediction with the test dataset.
6. Calculate performance metrics.

Code:

```
import numpy as np
import tensorflow as tf
from tensorflow import keras

from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import accuracy_score
# Generate a synthetic dataset

X, y = make_classification(n_samples=1000, n_features=20, random_state=42)
# Split the dataset into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
# Standardize features (optional but often beneficial)
scaler = StandardScaler()

X_train = scaler.fit_transform(X_train)
```

```

X_test = scaler.transform(X_test)
# Define the model
model = keras.Sequential([
    keras.layers.Input(shape=(X_train.shape[1],)), # Input layer
    keras.layers.Dense(64, activation='relu'), # Hidden layer with 64 neurons and ReLU activation
    keras.layers.Dense(1, activation='sigmoid') # Output layer with 1 neuron and sigmoid activation
])

# Train the model

history = model.fit(X_train, y_train, epochs=10, batch_size=32, validation_split=0.1)
# Evaluate the model on the test set
y_pred = model.predict(X_test)
y_pred_classes = (y_pred > 0.5).astype(int)
# Calculate accuracy on the test set
accuracy = accuracy_score(y_test, y_pred_classes)
# Calculate test loss
test_loss = model.evaluate(X_test, y_test)
print(f"Test accuracy: {accuracy * 100:.2f}%")
print(f"Test loss: {test_loss[0]:.4f}")

```

Output:

```

Epoch 1/10
192/192 [=====] - 5s 17ms/step - loss: 0.3739 - accuracy: 0.8950 - val_loss: 0.1801 - val_accuracy: 0.
9480
Epoch 2/10
192/192 [=====] - 3s 14ms/step - loss: 0.1492 - accuracy: 0.9562 - val_loss: 0.1261 - val_accuracy: 0.
9635
Epoch 3/10
192/192 [=====] - 2s 13ms/step - loss: 0.0980 - accuracy: 0.9714 - val_loss: 0.1129 - val_accuracy: 0.
9676
Epoch 4/10
192/192 [=====] - 2s 11ms/step - loss: 0.0711 - accuracy: 0.9795 - val_loss: 0.0962 - val_accuracy: 0.
9709
Epoch 5/10
192/192 [=====] - 2s 10ms/step - loss: 0.0543 - accuracy: 0.9844 - val_loss: 0.0914 - val_accuracy: 0.
9725
Epoch 6/10
192/192 [=====] - 2s 11ms/step - loss: 0.0402 - accuracy: 0.9888 - val_loss: 0.0866 - val_accuracy: 0.
9737
Epoch 7/10
192/192 [=====] - 2s 12ms/step - loss: 0.0301 - accuracy: 0.9920 - val_loss: 0.0871 - val_accuracy: 0.
9750
Epoch 8/10
192/192 [=====] - 2s 12ms/step - loss: 0.0245 - accuracy: 0.9931 - val_loss: 0.0840 - val_accuracy: 0.
9762
Epoch 9/10
192/192 [=====] - 2s 12ms/step - loss: 0.0180 - accuracy: 0.9956 - val_loss: 0.0878 - val_accuracy: 0.
9760
Epoch 10/10
192/192 [=====] - 2s 11ms/step - loss: 0.0149 - accuracy: 0.9963 - val_loss: 0.0858 - val_accuracy: 0.
9755

```

Jupyter Exp-1 Last Checkpoint: 3 hours ago (autosaved)

File Edit View Insert Cell Kernel Widgets Help Trusted Python 3 (ipykernel) Logout

```
In [1]: import numpy as np
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow.keras.datasets import mnist
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.utils import to_categorical

In [2]: feature_vector_length = 784
num_classes=10

In [3]: (X_train, Y_train), (X_test, Y_test) = mnist.load_data()


Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 [=====] - 2s 0us/step

In [4]: input_shape = (feature_vector_length)
print(f'Feature shape: {input_shape}')

Feature shape: 784

In [5]: X_train = X_train.reshape(X_train.shape[0], feature_vector_length)
X_test = X_test.reshape(X_test.shape[0], feature_vector_length)

In [6]: X_train = X_train.astype('float32') / 255
X_test = X_test.astype('float32') / 255
Y_train = to_categorical(Y_train, num_classes)
Y_test = to_categorical(Y_test, num_classes)
```



Result:

Thus, the implementation to build a simple neural network using Keras/TensorFlow has been successfully executed.

EX NO:2 **BUILD A CONVOLUTIONAL NEURAL NETWORK**
DATE:21/07/2025 **USING KERAS/TENSORFLOW**

Aim:

To implement a Convolutional Neural Network (CNN) using Keras/TensorFlow to recognize and classify handwritten digits from the MNIST dataset with high accuracy.

Procedure:

1. Import required libraries (TensorFlow/Keras, NumPy, etc.).
2. Load the MNIST dataset from Keras.
3. Normalize and reshape the image data.
4. Convert labels to one-hot encoded vectors.
5. Build a CNN model with Conv2D, MaxPooling, Flatten, and Dense layers.
6. Compile the model using categorical crossentropy and Adam optimizer.
7. Train the model on training data.
8. Evaluate the model on test data.
9. Display accuracy and predictions.

Code:

```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Dropout
from tensorflow.keras.datasets import mnist
import matplotlib.pyplot as plt
import numpy as np

(train_images, train_labels), (test_images, test_labels) = mnist.load_data()
train_images = train_images / 255.0
test_images = test_images / 255.0
train_images = train_images.reshape(-1, 28, 28, 1)
test_images = test_images.reshape(-1, 28, 28, 1)
model = Sequential([
    Conv2D(32, (3, 3), activation='relu', input_shape=(28, 28, 1)),
    MaxPooling2D((2, 2)),
    Conv2D(64, (3, 3), activation='relu'),
    MaxPooling2D((2, 2)),
    Flatten(),
    Dense(64, activation='relu'),
    Dropout(0.5),
    Dense(10, activation='softmax')
])
```

```

model.compile(optimizer='adam',
loss='sparse_categorical_crossentropy',
metrics=['accuracy'])

history = model.fit(train_images, train_labels,
epochs=5,
batch_size=64,
validation_split=0.2)

test_loss, test_acc = model.evaluate(test_images, test_labels)
print(f"\n Test accuracy: {test_acc:.4f}")
print(f" Test loss: {test_loss:.4f}")

plt.figure(figsize=(12, 5))
plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Accuracy', marker='o')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy', marker='o')
plt.title('Training and Validation Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()
plt.grid(True)

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss', marker='o')
plt.plot(history.history['val_loss'], label='Validation Loss', marker='o')
plt.title('Training and Validation Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.grid(True)
plt.tight_layout()
plt.show()

predictions = model.predict(test_images)
predicted_labels = np.argmax(predictions, axis=1)

num_samples = 10
plt.figure(figsize=(15, 4))

for i in range(num_samples):
    plt.subplot(1, num_samples, i + 1)
    plt.imshow(test_images[i].reshape(28, 28), cmap='gray')
    plt.title(f"Pred: {predicted_labels[i]}\nTrue: {test_labels[i]}")

```

```

plt.axis('off')
plt.suptitle("Sample Predictions on Test Images", fontsize=16)
plt.show()

```



Output:

```

Epoch 1/5
750/750 [=====] - 30s 39ms/step - loss: 0.3961 - accuracy: 0.8765 - val_loss: 0.0806 - val_accuracy: 0.9756
Epoch 2/5
750/750 [=====] - 26s 35ms/step - loss: 0.1538 - accuracy: 0.9548 - val_loss: 0.0606 - val_accuracy: 0.9824
Epoch 3/5
750/750 [=====] - 30s 39ms/step - loss: 0.1163 - accuracy: 0.9652 - val_loss: 0.0495 - val_accuracy: 0.9862
Epoch 4/5
750/750 [=====] - 27s 36ms/step - loss: 0.0937 - accuracy: 0.9725 - val_loss: 0.0472 - val_accuracy: 0.9872
Epoch 5/5
750/750 [=====] - 26s 35ms/step - loss: 0.0840 - accuracy: 0.9748 - val_loss: 0.0460 - val_accuracy: 0.9867
313/313 [=====] - 2s 5ms/step - loss: 0.0390 - accuracy: 0.9880

Test accuracy: 0.9880
Test loss: 0.0390

```



Result:

Thus, the Convolution Neural Network (CNN) using Keras / Tensorflow to recognize and classify handwritten digits from MNIST dataset has been implemented successfully.

EX NO: 3 IMAGE CLASSIFICATION ON CIFAR-10 DATASET USING CNN

DATE:28/07/2025

Aim:

To build a Convolutional Neural Network (CNN) model for classifying images from the CIFAR-10 dataset into one of the ten categories such as airplanes, cars, birds, cats, etc.

Procedure:

1. Download and load the CIFAR-10 dataset using Keras/TensorFlow.
2. Visualize and analyze sample images from the dataset.
3. Preprocess the data:
 - Normalize the pixel values (divide by 255)
 - Convert class labels to one-hot encoded format
4. Build a CNN model using Keras/TensorFlow:
 - Include convolutional, pooling, flatten, and dense layers.
5. Compile the model with suitable loss function and optimizer.
6. Train the model using training data and validate using test data.
7. Evaluate the model using accuracy and loss on test dataset.
8. Perform predictions on new/unseen CIFAR-10 images.
9. Visualize prediction results with sample images and predicted labels.

Code:

```
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
(x_train, y_train), (x_test, y_test) = tf.keras.datasets.cifar10.load_data()
x_train = x_train.astype('float32') / 255.0
x_test = x_test.astype('float32') / 255.0
y_train = tf.keras.utils.to_categorical(y_train, 10)
y_test = tf.keras.utils.to_categorical(y_test, 10)
model = tf.keras.Sequential()
model.add(tf.keras.layers.Conv2D(32, (3,3), activation='relu', input_shape=(32,32,3)))
model.add(tf.keras.layers.MaxPooling2D((2,2)))
model.add(tf.keras.layers.Conv2D(64, (3,3), activation='relu'))
model.add(tf.keras.layers.MaxPooling2D((2,2)))
model.add(tf.keras.layers.Conv2D(64, (3,3), activation='relu'))
model.add(tf.keras.layers.Flatten())
model.add(tf.keras.layers.Dense(64, activation='relu'))
model.add(tf.keras.layers.Dense(10, activation='softmax'))
model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])
model.fit(x_train, y_train, epochs=10, batch_size=64, validation_split=0.2)
class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
```

```

'dog', 'frog', 'horse', 'ship', 'truck']
index = int(input("Enter an index (0 to 9999) for test image: "))
if index < 0 or index >= len(x_test):
    print("Invalid index. Using index 0 by default.")
index = 0
test_image = x_test[index]
true_label = np.argmax(y_test[index])
prediction = model.predict(np.expand_dims(test_image, axis=0))
predicted_label = np.argmax(prediction)
plt.figure(figsize=(4, 4))
resized_image = tf.image.resize(test_image, [128, 128])
plt.imshow(resized_image)
plt.axis('off')
plt.title(f"Predicted: {class_names[predicted_label]}\nActual: {class_names[true_label]}")
plt.show()

```

Output:

```

Epoch 1/10
625/625 [=====] - 58s 87ms/step - loss: 1.6801 - accuracy: 0.3846 - val_loss: 1.4341 - val_accuracy:
0.4803
Epoch 2/10
625/625 [=====] - 37s 60ms/step - loss: 1.3153 - accuracy: 0.5284 - val_loss: 1.3005 - val_accuracy:
0.5388
Epoch 3/10
625/625 [=====] - 36s 58ms/step - loss: 1.1663 - accuracy: 0.5846 - val_loss: 1.1370 - val_accuracy:
0.6014
Epoch 4/10
625/625 [=====] - 38s 61ms/step - loss: 1.0629 - accuracy: 0.6249 - val_loss: 1.0984 - val_accuracy:
0.6178
Epoch 5/10
625/625 [=====] - 41s 65ms/step - loss: 0.9991 - accuracy: 0.6480 - val_loss: 1.0476 - val_accuracy:
0.6379
Epoch 6/10
625/625 [=====] - 38s 61ms/step - loss: 0.9348 - accuracy: 0.6720 - val_loss: 0.9795 - val_accuracy:
0.6598
Epoch 7/10
625/625 [=====] - 38s 60ms/step - loss: 0.8764 - accuracy: 0.6970 - val_loss: 1.0013 - val_accuracy:
0.6547
Epoch 8/10
625/625 [=====] - 38s 61ms/step - loss: 0.8338 - accuracy: 0.7096 - val_loss: 0.9313 - val_accuracy:
0.6770
Epoch 9/10
625/625 [=====] - 39s 62ms/step - loss: 0.7943 - accuracy: 0.7242 - val_loss: 0.9243 - val_accuracy:
0.6856
Epoch 10/10
625/625 [=====] - 37s 60ms/step - loss: 0.7588 - accuracy: 0.7362 - val_loss: 0.8994 - val_accuracy:
0.6986

```

Predicted: frog
 Actual: frog



Result:

Thus, the Convolution Neural Network (CNN) model for classifying images from CIFAR-10 dataset is implemented successfully.

Ex No: 4 TRANSFER LEARNING WITH CNN AND VISUALIZATION

DATE:04/08/2025

Aim:

To build a convolutional neural network with transfer learning and perform visualization

Procedure:

1. Download and load the dataset.
2. Perform analysis and preprocessing of the dataset.
3. Build a simple neural network model using Keras/TensorFlow.
4. Compile and fit the model.
5. Perform prediction with the test dataset.
6. Calculate performance metrics.

Code:

```
conda install -c conda-forge python-graphviz -y
import tensorflow as tf
from tensorflow.keras.applications import VGG16
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Flatten, Dropout
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.datasets import cifar10
from tensorflow.keras.utils import plot_model
import matplotlib.pyplot as plt
import numpy as np
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
x_train = x_train / 255.0
x_test = x_test / 255.0
vgg_base = VGG16(weights='imagenet', include_top=False, input_shape=(32, 32, 3))
for layer in vgg_base.layers:
    layer.trainable = False
model = Sequential()
model.add(vgg_base)
model.add(Flatten())
model.add(Dense(512, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(10, activation='softmax'))
model.compile(optimizer=Adam(learning_rate=0.0001),
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
plot_model(model, to_file='cnn.png', show_shapes=True,
```

```
show_layer_names=True, dpi=300)
plt.figure(figsize=(20, 20))
img = plt.imread('cnn.png')
plt.imshow(img)
plt.axis('off')
plt.show()
history = model.fit(x_train, y_train,
epochs=10,
batch_size=32,
validation_split=0.2)

test_loss, test_acc = model.evaluate(x_test, y_test)
print(f'Test Loss: {test_loss:.4f}')
print(f'Test Accuracy: {test_acc * 100:.2f}%')
plt.figure(figsize=(12, 5))


plt.subplot(1, 2, 1)
plt.plot(history.history['accuracy'], label='Train Accuracy')
plt.plot(history.history['val_accuracy'], label='Validation Accuracy')
plt.title('Model Accuracy')
plt.xlabel('Epoch')
plt.ylabel('Accuracy')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(history.history['loss'], label='Train Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.title('Model Loss')
plt.xlabel('Epoch')
plt.ylabel('Loss')
plt.legend()
plt.tight_layout()
plt.show()

class_names = ['airplane', 'automobile', 'bird', 'cat', 'deer',
'dog', 'frog', 'horse', 'ship', 'truck']
sample = x_test[0].reshape(1, 32, 32, 3)
prediction = model.predict(sample)
predicted_class = class_names[np.argmax(prediction)]

plt.imshow(x_test[0])
plt.title(f'Predicted: {predicted_class}')
plt.axis('off')
plt.show()
```


Output:



```

Epoch 1/10
1250/1250 [=====] - 231s 182ms/step - loss: 1.7748 - accuracy: 0.3727 - val_loss: 1.4674 - val_accuracy: 0.4959
Epoch 2/10
1250/1250 [=====] - 193s 154ms/step - loss: 1.4665 - accuracy: 0.4920 - val_loss: 1.3556 - val_accuracy: 0.5322
Epoch 3/10
1250/1250 [=====] - 187s 150ms/step - loss: 1.3733 - accuracy: 0.5264 - val_loss: 1.2966 - val_accuracy: 0.5512
Epoch 4/10
1250/1250 [=====] - 189s 151ms/step - loss: 1.3197 - accuracy: 0.5386 - val_loss: 1.2610 - val_accuracy: 0.5621
Epoch 5/10
1250/1250 [=====] - 191s 153ms/step - loss: 1.2777 - accuracy: 0.5551 - val_loss: 1.2352 - val_accuracy: 0.5739
Epoch 6/10
1250/1250 [=====] - 190s 152ms/step - loss: 1.2474 - accuracy: 0.5683 - val_loss: 1.2154 - val_accuracy: 0.5759
Epoch 7/10
1250/1250 [=====] - 187s 150ms/step - loss: 1.2269 - accuracy: 0.5741 - val_loss: 1.1981 - val_accuracy: 0.5830
Epoch 8/10
1250/1250 [=====] - 183s 146ms/step - loss: 1.2039 - accuracy: 0.5834 - val_loss: 1.1839 - val_accuracy: 0.5864
Epoch 9/10
1250/1250 [=====] - 177s 142ms/step - loss: 1.1866 - accuracy: 0.5887 - val_loss: 1.1735 - val_accuracy: 0.5900
Epoch 10/10
1250/1250 [=====] - 175s 140ms/step - loss: 1.1699 - accuracy: 0.5943 - val_loss: 1.1672 - val_accuracy: 0.5910

```



Result:

Thus, the Convolution Neural Network (CNN) with transfer learning and perform visualization has been implemented successfully

Aim:

To build a recurrent neural network with Keras/TensorFlow.

Procedure:

1. Download and load the dataset.
 2. Perform analysis and preprocessing of the dataset.
 3. Build a simple neural network model using Keras/TensorFlow.
 4. Compile and fit the model.
 5. Perform prediction with the test dataset.
 6. Calculate performance metrics.

Code:

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import SimpleRNN, Dense
from sklearn.metrics import r2_score
np.random.seed(0)
seq_length = 10
num_samples = 1000
X = np.random.randn(num_samples, seq_length, 1)
y = X.sum(axis=1) + 0.1 * np.random.randn(num_samples, 1)
split_ratio = 0.8
split_index = int(split_ratio * num_samples)
X_train, X_test = X[:split_index], X[split_index:]
y_train, y_test = y[:split_index], y[split_index:]
model = Sequential()
model.add(SimpleRNN(units=50, activation='relu', input_shape=(seq_length, 1)))
model.add(Dense(units=1))
model.compile(optimizer='adam', loss='mean_squared_error')
model.summary()
batch_size = 30
epochs = 50 # Reduced epochs for quick demonstration
history = model.fit(
    X_train, y_train,
    batch_size=batch_size,
    epochs=epochs,
    validation_split=0.2
)
test_loss = model.evaluate(X_test, y_test)
print(f'Test Loss: {test_loss:.4f}')
```

```

y_pred = model.predict(X_test)
r2 = r2_score(y_test, y_pred)
print(f'Test Accuracy (R^2): {r2:.4f}')

new_data = np.random.randn(5, seq_length, 1)
predictions = model.predict(new_data)
print("Predictions for new data:")
print(predictions)

```

Output:

Model: "sequential"

Layer (type)	Output Shape	Param #
simple_rnn (SimpleRNN)	(None, 50)	2600
dense (Dense)	(None, 1)	51
<hr/>		
Total params: 2,651		
Trainable params: 2,651		
Non-trainable params: 0		

```

Epoch 1/50
22/22 [=====] - 2s 23ms/step - loss: 8.7454
- val_loss: 6.3263
Epoch 2/50
22/22 [=====] - 0s 4ms/step - loss: 5.8837
- val_loss: 3.7798
Epoch 3/50
22/22 [=====] - 0s 5ms/step - loss: 3.7728
- val_loss: 2.3105
Epoch 4/50
22/22 [=====] - 0s 5ms/step - loss: 1.7141
- val_loss: 0.5373
Epoch 5/50
22/22 [=====] - 0s 4ms/step - loss: 0.2878
- val_loss: 0.2417
Epoch 6/50
22/22 [=====] - 0s 4ms/step - loss: 0.1304
- val_loss: 0.1146
Epoch 7/50

1/1 [=====] - 0s 20ms/step
Predictions for new data:
[[ 1.5437698]
 [ 0.4290885]
 [-2.1180325]
 [-0.5443404]
 [-3.8416493]]

```

Result:

Thus, the Recurrent Neural Network (RNN) has been implemented using Tensorflow.

EX NO: 6 SENTIMENT CLASSIFICATION OF TEXT USING RNN**DATE:15/09/2025****Aim:**

To implement a Recurrent Neural Network (RNN) using Keras/TensorFlow for classifying the sentiment of text data (e.g., movie reviews) as positive or negative.

Procedure:

1. Import necessary libraries.
2. Load and preprocess the text dataset (e.g., IMDb).
3. Pad sequences and prepare labels.
4. Build an RNN model with Embedding and SimpleRNN layers.
5. Compile the model with loss and optimizer.
6. Train the model on training data.
7. Evaluate the model on test data.
8. Predict sentiment for new inputs

Code:

```
import numpy as np
from tensorflow.keras.datasets import imdb
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, SimpleRNN, Dense
max_words = 5000
max_len = 200
(x_train, y_train), (x_test, y_test) = imdb.load_data(num_words=max_words)
X_train = pad_sequences(x_train, maxlen=max_len)
X_test = pad_sequences(x_test, maxlen=max_len)
model = Sequential()
model.add(Embedding(input_dim=max_words, output_dim=32, input_length=max_len))
model.add(SimpleRNN(32))
model.add(Dense(1, activation='sigmoid'))
model.compile(optimizer='adam', loss='binary_crossentropy', metrics=['accuracy'])
print("Training...")
model.fit(X_train, y_train, epochs=2, batch_size=64, validation_split=0.2)
loss, acc = model.evaluate(X_test, y_test)
print(f"\nTest Accuracy: {acc:.4f}")
word_index = imdb.get_word_index()
reverse_word_index = {v: k for (k, v) in word_index.items()}
```

```

def decode_review(review):
    return " ".join([reverse_word_index.get(i - 3, "?") for i in review])
sample_review = X_test[0]
prediction = model.predict(sample_review.reshape(1, -1))[0][0]
print("\nReview text:", decode_review(x_test[0]))
print("Predicted Sentiment:", "Positive " if prediction > 0.5 else "Negative ")

```

Output:

```

→ Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb.npz
17464789/17464789 ━━━━━━━━ 0s 0us/step
Training...
Epoch 1/2
/usr/local/lib/python3.12/dist-packages/keras/src/layers/core/embedding.py:97: UserWarning: Argument `input_length` is deprecated. Just remove it.
    warnings.warn(
313/313 ━━━━━━━━ 21s 59ms/step - accuracy: 0.6479 - loss: 0.6143 - val_accuracy: 0.6644 - val_loss: 0.6085
Epoch 2/2
313/313 ━━━━━━━━ 17s 53ms/step - accuracy: 0.7939 - loss: 0.4496 - val_accuracy: 0.8186 - val_loss: 0.4121
782/782 ━━━━━━ 10s 13ms/step - accuracy: 0.8237 - loss: 0.4115

Test Accuracy: 0.8230
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/imdb\_word\_index.json
1641221/1641221 ━━━━━━━━ 0s 0us/step

```

```

● def decode_review(review):
    return " ".join([reverse_word_index.get(i - 3, "?") for i in review])
sample_review = X_test[0]
prediction = model.predict(sample_review.reshape(1, -1))[0][0]
print("\nReview text:", decode_review(x_test[0]))
print("Predicted Sentiment:", "Positive " if prediction > 0.5 else "Negative ")

```

→ 1/1 ━━━ 0s 195ms/step

Review text: ? please give this one a miss br br ? ? and the rest of the cast ? terrible performances the show is flat flat flat br br i don't know how michael
Predicted Sentiment: Negative

Result:

Thus, the Recurrent Neural Network (RNN) using Keras has been implemented for classifying sentiment of text successfully.

Ex No: 7 BUILD AUTOENCODERS WITH KERAS/TENSORFLOW**DATE:22/09/2025****Aim:**

To build autoencoders with Keras/TensorFlow.

Procedure:

1. Download and load the dataset.
2. Perform analysis and preprocessing of the dataset.
3. Build a simple neural network model using Keras/TensorFlow.
4. Compile and fit the model.
5. Perform prediction with the test dataset.
6. Calculate performance metrics.

Code:

```
import numpy as np
import matplotlib.pyplot as plt
from keras.layers import Input, Dense
from keras.models import Model
from keras.datasets import mnist
(x_train, _), (x_test, _) = mnist.load_data()
x_train = x_train.astype('float32') / 255.
x_test = x_test.astype('float32') / 255.
x_train = x_train.reshape((len(x_train), np.prod(x_train.shape[1:])))
x_test = x_test.reshape((len(x_test), np.prod(x_test.shape[1:])))
input_img = Input(shape=(784,))
encoded = Dense(32, activation='relu')(input_img)
decoded = Dense(784, activation='sigmoid')(encoded)
autoencoder = Model(input_img, decoded)
autoencoder.compile(optimizer='adam', loss='binary_crossentropy')
autoencoder.fit(x_train, x_train,
                epochs=50,
                batch_size=256,
                shuffle=True,
                validation_data=(x_test, x_test))
test_loss = autoencoder.evaluate(x_test, x_test)
decoded_imgs = autoencoder.predict(x_test)
threshold = 0.5
correct_predictions = np.sum(
    np.where(x_test >= threshold, 1, 0) ==
    np.where(decoded_imgs >= threshold, 1, 0))
total_pixels = x_test.shape[0] * x_test.shape[1]
test_accuracy = correct_predictions / total_pixels
print("Test Loss:", test_loss)
```

```

print("Test Accuracy:", test_accuracy)
n = 10
plt.figure(figsize=(20, 4))
for i in range(n):
    # Display original
    ax = plt.subplot(2, n, i + 1)
    plt.imshow(x_test[i].reshape(28, 28))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)

    # Display reconstruction with threshold
    ax = plt.subplot(2, n, i + 1 + n)
    reconstruction = decoded_imgs[i].reshape(28, 28)
    plt.imshow(np.where(reconstruction >= threshold, 1.0, 0.0))
    plt.gray()
    ax.get_xaxis().set_visible(False)
    ax.get_yaxis().set_visible(False)
plt.show()

```

Output:

```

Epoch 1/50
235/235 ━━━━━━━━ 6s 18ms/step - loss: 0.3805 - val_loss: 0.1906
Epoch 2/50
235/235 ━━━━━━ 5s 19ms/step - loss: 0.1808 - val_loss: 0.1547
Epoch 3/50
235/235 ━━━━ 5s 19ms/step - loss: 0.1501 - val_loss: 0.1342
Epoch 4/50
235/235 ━━━━ 3s 10ms/step - loss: 0.1321 - val_loss: 0.1221
Epoch 5/50
235/235 ━━━━ 2s 9ms/step - loss: 0.1210 - val_loss: 0.1138
Epoch 6/50
235/235 ━━━━ 3s 11ms/step - loss: 0.1134 - val_loss: 0.1081
Epoch 7/50
235/235 ━━━━ 5s 9ms/step - loss: 0.1079 - val_loss: 0.1039
Epoch 8/50
235/235 ━━━━ 2s 9ms/step - loss: 0.1042 - val_loss: 0.1006
Epoch 9/50
235/235 ━━━━ 3s 9ms/step - loss: 0.1011 - val_loss: 0.0981
Epoch 10/50
235/235 ━━━━ 3s 11ms/step - loss: 0.0989 - val_loss: 0.0963
Epoch 11/50
235/235 ━━━━ 3s 12ms/step - loss: 0.0972 - val_loss: 0.0951
Epoch 12/50
235/235 ━━━━ 3s 11ms/step - loss: 0.0964 - val_loss: 0.0943
Epoch 13/50
235/235 ━━━━ 2s 10ms/step - loss: 0.0954 - val_loss: 0.0938
Epoch 14/50
235/235 ━━━━ 2s 10ms/step - loss: 0.0950 - val_loss: 0.0934
Epoch 15/50
235/235 ━━━━ 3s 11ms/step - loss: 0.0944 - val_loss: 0.0932

```

```
→ Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11490434/11490434 1s 0us/step
```

```
→ Test Loss: 0.09166844934225082
Test Accuracy: 0.9712756377551021
```



```
# Display reconstruction with threshold
ax = plt.subplot(2, n, i + 1 + n)
reconstruction = decoded_imgs[i].reshape(28, 28)
plt.imshow(np.where(reconstruction >= threshold, 1.0, 0.0))
plt.gray()
ax.get_xaxis().set_visible(False)
ax.get_yaxis().set_visible(False)
plt.show()
```



Result:

Thus, an Autoencoder has been implemented using Keras / Tensorflow.

Ex No:8

OBJECT DETECTION WITH YOLO3

DATE:29/09/2025

Aim:

To build an object detection model with YOLO3 using Keras/TensorFlow.

Procedure:

1. Download and load the dataset.
2. Perform analysis and preprocessing of the dataset.
3. Build a simple neural network model using Keras/TensorFlow.
4. Compile and fit the model.
5. Perform prediction with the test dataset.
6. Calculate performance metrics.

Code:

```
import cv2
import matplotlib.pyplot as plt
import numpy as np

# Define the paths to the YOLOv3 configuration, weights, and class names files
cfg_file = '/content/yolov3.cfg'
weight_file = '/content/yolov3.weights'
namesfile = '/content/coco.names'

# Load the YOLOv3 model
net = cv2.dnn.readNet(weight_file, cfg_file)

# Load class names
with open(namesfile, 'r') as f:
    classes = f.read().strip().split('\n')

# Load an image for object detection
image_path = '/content/hit.jpg'
image = cv2.imread(image_path)

# Get the height and width of the image
height, width = image.shape[:2]

# Create a blob from the image
blob = cv2.dnn.blobFromImage(image, 1/255.0, (416, 416), swapRB=True, crop=False)
net.setInput(blob)

# Get the names of the output layers
layer_names = net.getUnconnectedOutLayersNames()

# Run forward pass
```

```

outs = net.forward(layer_names)

# Initialize lists to store detected objects' information
class_ids = []
confidences = []
boxes = []

# Define a confidence threshold for object detection
conf_threshold = 0.5

# Loop over the detections
for out in outs:
    for detection in out:
        scores = detection[5:]
        class_id = np.argmax(scores)
        confidence = scores[class_id]
        if confidence > conf_threshold:
            # Object detected
            center_x = int(detection[0] * width)
            center_y = int(detection[1] * height)
            w = int(detection[2] * width)
            h = int(detection[3] * height)

            # Rectangle coordinates
            x = int(center_x - w / 2)
            y = int(center_y - h / 2)

            class_ids.append(class_id)
            confidences.append(float(confidence))
            boxes.append([x, y, w, h])

# Apply non-maximum suppression to eliminate overlapping boxes
nms_threshold = 0.4
indices = cv2.dnn.NMSBoxes(boxes, confidences, conf_threshold, nms_threshold)

# Draw bounding boxes and labels on the image
for i in indices.flatten(): # flatten for compatibility
    x, y, w, h = boxes[i]
    label = str(classes[class_ids[i]])
    confidence = confidences[i]

    cv2.rectangle(image, (x, y), (x + w, y + h), (0, 255, 0), 2)
    cv2.putText(image, f'{label} {confidence:.2f}', (x, y - 10),
               cv2.FONT_HERSHEY_SIMPLEX, 0.8, (0, 255, 0), 2)

```

```
# Display the result in Jupyter Notebook
plt.figure(figsize=(10, 8))
plt.imshow(cv2.cvtColor(image, cv2.COLOR_BGR2RGB))
plt.axis('off')
plt.show()
```

Output:


```
cfg_file=r"C:\Users\agaki\Downloads\yolov3.cfg"
weight_file=r"C:\Users\agaki\Downloads\yolov3.weights"
names_file=r"C:\Users\agaki\Downloads\coco.names"

for out in outs:
    for detection in out:
        scores=detection[5:]
        class_id=np.argmax(scores)
        confidence=scores[class_id]
        if confidence>conf_threshold:

            center_x=int(detection[0]*width)
            center_y=int(detection[1]*height)
            w=int(detection[2]*width)
            h=int(detection[3]*height)

            x=int(center_x-w/2)
            y=int(center_y-h/2)

            class_ids.append(class_id)
            confidences.append(float(confidence))
            boxes.append([x,y,w,h])
```



Result:

Thus, object detection using YOLOV5 has been implemented successfully.

Ex No: 9 BUILD GENERATIVE ADVERSARIAL NEURAL NETWORK**DATE:29/09/2025****Aim:**

To build a generative adversarial neural network using Keras/TensorFlow.

Procedure:

1. Download and load the dataset.
2. Perform analysis and preprocessing of the dataset.
3. Build a simple neural network model using Keras/TensorFlow.
4. Compile and fit the model.
5. Perform prediction with the test dataset.
6. Calculate performance metrics.

Code:

```
import numpy as np
import tensorflow as tf
from tensorflow.keras.layers import Dense
from tensorflow.keras.models import Sequential
from tensorflow.keras.optimizers import Adam
from sklearn.datasets import load_iris
import matplotlib.pyplot as plt

# Load and Preprocess the Iris Dataset
iris = load_iris()
x_train = iris.data

# Build the GAN model
def build_generator():
    model = Sequential()
    model.add(Dense(128, input_shape=(100,), activation='relu'))
    model.add(Dense(4, activation='linear')) # Output 4 features
    return model

def build_discriminator():
    model = Sequential()
    model.add(Dense(128, input_shape=(4,), activation='relu'))
    model.add(Dense(1, activation='sigmoid'))
    return model

def build_gan(generator, discriminator):
    discriminator.trainable = False
    model = Sequential()
    model.add(generator)
    model.add(discriminator)
```

```

return model

generator = build_generator()
discriminator = build_discriminator()
gan = build_gan(generator, discriminator)

# Compile the Models
generator.compile(loss='mean_squared_error', optimizer=Adam(0.0002, 0.5))
discriminator.compile(loss='binary_crossentropy', optimizer=Adam(0.0002, 0.5),
metrics=['accuracy'])
gan.compile(loss='binary_crossentropy', optimizer=Adam(0.0002, 0.5))

# Training Loop
epochs = 200
batch_size = 16

for epoch in range(epochs):
    # Train discriminator
    idx = np.random.randint(0, x_train.shape[0], batch_size)
    real_samples = x_train[idx]
    fake_samples = generator.predict(np.random.normal(0, 1, (batch_size, 100)), verbose=0)

    real_labels = np.ones((batch_size, 1))
    fake_labels = np.zeros((batch_size, 1))

    d_loss_real = discriminator.train_on_batch(real_samples, real_labels)
    d_loss_fake = discriminator.train_on_batch(fake_samples, fake_labels)

    # Train generator
    noise = np.random.normal(0, 1, (batch_size, 100))
    g_loss = gan.train_on_batch(noise, real_labels)

    # Print progress
    print(f"Epoch {epoch}/{epochs} | Discriminator Loss: {0.5 * (d_loss_real[0] + d_loss_fake[0])} | Generator Loss: {g_loss}")

    # Generating Synthetic Data
    synthetic_data = generator.predict(np.random.normal(0, 1, (150, 100)), verbose=0)

    # Create scatter plots for feature pairs
    plt.figure(figsize=(12, 8))
    plot_idx = 1

    for i in range(4):
        for j in range(i + 1, 4):
            plt.subplot(2, 3, plot_idx)

```

```

plt.scatter(x_train[:, i], x_train[:, j], label='Real Data', c='blue', marker='o', s=30)
plt.scatter(synthetic_data[:, i], synthetic_data[:, j], label='Synthetic Data', c='red', marker='x',
s=30)
plt.xlabel(f'Feature {i + 1}')
plt.ylabel(f'Feature {j + 1}')
plt.legend()
plot_idx += 1

plt.tight_layout()
plt.show()

```

Output:

```

Epoch 0/200 | Discriminator Loss: 0.8773080408573151 |Generator Loss: 0.764731228351593
Epoch 1/200 | Discriminator Loss: 0.9332943856716156 |Generator Loss: 0.7988691329956055
Epoch 2/200 | Discriminator Loss: 0.9277275502681732 |Generator Loss: 0.8127573728561401
Epoch 3/200 | Discriminator Loss: 0.8921994566917419 |Generator Loss: 0.7757299542427963
Epoch 4/200 | Discriminator Loss: 0.913447916507721 |Generator Loss: 0.7737997174263
Epoch 5/200 | Discriminator Loss: 0.8916181325912476 |Generator Loss: 0.8003895282745361
Epoch 6/200 | Discriminator Loss: 0.9026078879833221 |Generator Loss: 0.814433217048645
Epoch 7/200 | Discriminator Loss: 0.9135120809078217 |Generator Loss: 0.8237183690071106
Epoch 8/200 | Discriminator Loss: 0.879832923412323 |Generator Loss: 0.7563657760620117
Epoch 9/200 | Discriminator Loss: 0.9439513385295868 |Generator Loss: 0.7623365521430969
Epoch 10/200 | Discriminator Loss: 0.9355685114860535 |Generator Loss: 0.7924684286117554
Epoch 11/200 | Discriminator Loss: 0.9386743983160095 |Generator Loss: 0.7614541053771973
Epoch 12/200 | Discriminator Loss: 0.96055225610733 |Generator Loss: 0.7792538404464722
Epoch 13/200 | Discriminator Loss: 0.9134297668933868 |Generator Loss: 0.792992115020752
Epoch 14/200 | Discriminator Loss: 0.8851655125617981 |Generator Loss: 0.7628173232078552
Epoch 15/200 | Discriminator Loss: 0.9505723416805267 |Generator Loss: 0.7851851582527161
Epoch 16/200 | Discriminator Loss: 0.92226842045784 |Generator Loss: 0.769191563129425
Epoch 17/200 | Discriminator Loss: 0.8982412815093994 |Generator Loss: 0.7685977220535278
Epoch 18/200 | Discriminator Loss: 0.9125983119010925 |Generator Loss: 0.7730982899665833
Epoch 19/200 | Discriminator Loss: 0.9367325385938721 |Generator Loss: 0.7837406396865845
Epoch 20/200 | Discriminator Loss: 0.9531015455722809 |Generator Loss: 0.7827053070068359
Epoch 21/200 | Discriminator Loss: 0.9306998252868652 |Generator Loss: 0.76667914032936096
Epoch 22/200 | Discriminator Loss: 0.8887360095977783 |Generator Loss: 0.7845874428749084
Epoch 23/200 | Discriminator Loss: 0.9426513016223907 |Generator Loss: 0.746765673160553
Epoch 24/200 | Discriminator Loss: 0.9331325888633728 |Generator Loss: 0.761589765548786
Epoch 25/200 | Discriminator Loss: 0.9080778360366821 |Generator Loss: 0.7709233164787292
Epoch 26/200 | Discriminator Loss: 0.9232879281044006 |Generator Loss: 0.7773635387420654
Epoch 27/200 | Discriminator Loss: 0.9102294743061066 |Generator Loss: 0.7809370756149292
Epoch 28/200 | Discriminator Loss: 0.9312145709991455 |Generator Loss: 0.7647197246551514
Epoch 29/200 | Discriminator Loss: 0.9415165781974792 |Generator Loss: 0.7561923861503601
Epoch 30/200 | Discriminator Loss: 0.930676281452179 |Generator Loss: 0.7709008455276489
Epoch 31/200 | Discriminator Loss: 0.9495892226696014 |Generator Loss: 0.7595088481903076

```


```


In [33]: synthetic_data = generator.predict(np.random.normal(0,1,(150,100)),verbose=0)
plt.figure(figsize=(12,8))
plot_idx1

for i in range(4):
    for j in range(i+1,4):
        plt.subplot(2,i,plot_idx)
        plt.scatter(x_train[:, i],x_train[:, j],label='Real Data',c='blue',marker='o',s=30)
        plt.scatter(synthetic_data[:, i],synthetic_data[:, j],label='Synthetic Data',c='red',marker='x',s=30)
        plt.xlabel(f'Feature{i+1}')
        plt.ylabel(f'Feature{j+1}')
        plt.legend()
        plot_idx+=1

    plt.tight_layout()
    plt.show()

```





Result:

Thus, a generative adversarial neural network using Keras / Tensorflow has been implemented successfully.

EX NO:10

DATE:07-10-2025

Mini Project: News Headline Prediction

Aim:

To develop a News Headline Generation model using Deep Learning (LSTM network) that can automatically generate meaningful headlines from textual news data of the Reuters dataset, thereby demonstrating the ability of machines to understand and summarize human-written information.

Algorithm:

1. Import libraries and load the Reuters dataset.
2. Decode sequences and clean the text data.
3. Tokenize the text and create n-gram sequences.
4. Pad sequences to equal length for model input.
5. Build a simple LSTM model with embedding and dense layers.
6. Compile the model using sparse categorical crossentropy and Adam optimizer.
7. Train the model and monitor validation performance.
8. Generate news headlines from seed text and evaluate the outputs.

Code:

```
import numpy as np import re
import matplotlib.pyplot as plt
from tensorflow.keras.datasets import reuters
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Embedding, LSTM, Dense
from tensorflow.keras.callbacks import EarlyStopping

# -----
# Parameters

# -----
NUM_WORDS = 10000    # Limit vocabulary SEQ_LEN = 10 # Shorter context -> faster
EMBED_DIM = 64
BATCH_SIZE = 64       # Smaller batch = faster iteration EPOCHS = 5      # Lower for
speed MAX_HEADLINE_WORDS = 6
```

```
# 1) Load dataset (Corrected unpacking)
(x_train, _), _ = reuters.load_data(num_words=NUM_WORDS)

# ✓ Reduce dataset size to 20% for speed x_train = x_train[:int(len(x_train) * 0.2)]

# 2) Word index mapping
raw_word_index = reuters.get_word_index()
index_to_word = {v+3: k for k,v in raw_word_index.items()} index_to_word[0] = '<PAD>'
index_to_word[1] = '<START>' index_to_word[2] = '<UNK>'

def decode_seq(seq):
    return ''.join([index_to_word.get(i,"") for i in seq])

texts = [decode_seq(seq) for seq in x_train]
texts = [re.sub(r'\d+', " ", t).lower() for t in texts] # remove numbers

# 3) Tokenize
tokenizer = Tokenizer(num_words=NUM_WORDS, oov_token='<UNK>')
```

```
tokenizer.fit_on_texts(texts)
vocab_size = min(NUM_WORDS, len(tokenizer.word_index)+1)

# 4) Build sequences all_sequences = [] for txt in texts:
token_list = tokenizer.texts_to_sequences([txt])[0] for i in range(1,len(token_list)):
start = max(0,i-SEQ_LEN+1) ngram = token_list[start:i+1] all_sequences.append(ngram)

padded = pad_sequences(all_sequences, maxlen=SEQ_LEN, padding='pre') X = padded[:, :-1]
y = padded[:, -1]

# 5) Fast Simple LSTM model model = Sequential([
Embedding(input_dim=vocab_size, output_dim=EMBED_DIM), LSTM(64),
Dense(vocab_size, activation='softmax')
])

model.compile(loss='sparse_categorical_crossentropy', optimizer='adam',
metrics=['sparse_categorical_accuracy'])
model.summary()

es = EarlyStopping(monitor='val_loss', patience=2, restore_best_weights=True)
history = model.fit(X, y, epochs=EPOCHS, batch_size=BATCH_SIZE, validation_split=0.1,
callbacks=[es])
```

```
# 6) Headline generator
def sample_preds(preds, temperature=1.0): preds = np.asarray(preds).astype('float64') preds =
np.log(preds+1e-12)/temperature preds = np.exp(preds)/np.sum(np.exp(preds))
return np.argmax(np.random.multinomial(1,preds,1))
```

```
def generate_headline(seed_text, max_words=6, temperature=0.7): seed_text = re.sub(r"[^a-
-zA-Z0-9\s]","",seed_text).lower() seed_seq = tokenizer.texts_to_sequences([seed_text])[0]
generated = []
for _ in range(max_words):
    input_seq = seed_seq[-(SEQ_LEN-1):] if len(seed_seq)>=SEQ_LEN-1 else seed_seq
    input_padded = pad_sequences([input_seq], maxlen=SEQ_LEN-1, padding='pre') preds =
model.predict(input_padded,verbose=0)[0]
    next_idx = sample_preds(preds,temperature) next_word =
tokenizer.index_word.get(next_idx,"") if next_word in [",<PAD>','<UNK>"]: break
    generated.append(next_word) seed_seq.append(next_idx)
return ''.join(generated).title()
```

```
# 7) Final Output (Like friend's format)
print("\n--- Sample News & Predicted Headlines ---") for i in range(5):
news = texts[np.random.randint(0,len(texts))]
headline = generate_headline(news, max_words=6, temperature=0.7) print(f"\nNews:
{news[:120]}...")
print(f'Predicted Headline: {headline}')
```

Output:

```
Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/reuters_word_index.json
550378/550378 - 0s 0us/step
Model: "sequential"
Layer (type)          Output Shape       Param #
embedding (Embedding)    ?           0 (unbuilt)
lstm (LSTM)            ?           0 (unbuilt)
dense (Dense)          ?           0 (unbuilt)
Total params: 0 (0.00 B)
Trainable params: 0 (0.00 B)
Non-trainable params: 0 (0.00 B)
Epoch 1/5
3367/3367 --> 110s 32ms/step - loss: 6.6616 - sparse_categorical_accuracy: 0.0696 - val_loss: 6.0166 - val_sparse_categorical_accuracy: 0.1260
Epoch 2/5
3367/3367 --> 112s 33ms/step - loss: 5.6497 - sparse_categorical_accuracy: 0.1499 - val_loss: 5.6980 - val_sparse_categorical_accuracy: 0.1556
Epoch 3/5
3367/3367 --> 137s 32ms/step - loss: 5.2343 - sparse_categorical_accuracy: 0.1845 - val_loss: 5.5349 - val_sparse_categorical_accuracy: 0.1781
Epoch 4/5
3367/3367 --> 108s 32ms/step - loss: 4.9603 - sparse_categorical_accuracy: 0.2061 - val_loss: 5.4541 - val_sparse_categorical_accuracy: 0.1872
Epoch 5/5
3367/3367 --> 141s 32ms/step - loss: 4.7516 - sparse_categorical_accuracy: 0.2212 - val_loss: 5.4107 - val_sparse_categorical_accuracy: 0.1969
--- Sample News & Predicted Headlines ---
News: <start> oper shr loss cts vs profit three cts oper net loss vs profit revs vs mln note year ago oper net exclu...
Predicted Headline: Dlrs Vs Reuter Gains Net Mln
News: <start> net vs year net vs note company's full name is indiana federal savings and loan association per share ...
Predicted Headline: From Mln Dlrs In The Refiners
News: <start> the u s economy is showing some promising signs of accelerated expansion despite the <unk> of the fourth quarter...
Predicted Headline: For The First Season In The
News: <start> japan's seasonally adjusted unemployment rate rose to a record pct in january the worst since the government s...
Predicted Headline: Gillette Is Unk We Have To
```

Result:

Thus, my project generates news headlines from text using an LSTM-based deep learning model and the output is verified successfully